

**NAME**

**mac, GLDv3** - MAC networking device driver overview

**SYNOPSIS**

```
#include <sys/mac_provider.h>
#include <sys/mac_ether.h>
```

**INTERFACE LEVEL**

illumos DDI specific

**DESCRIPTION**

The **MAC** framework provides a means for implementing high-performance networking device drivers. It is the successor to the GLD interfaces and is sometimes referred to as the GLDv3. The remainder of this manual introduces the aspects of writing devices drivers that leverage the MAC framework. While both the GLDv3 and MAC framework refer to the same thing, in this manual page we use the term the *MAC framework* to refer to the device driver interface.

MAC device drivers are character devices. They define the standard `_init(9E)`, `_fini(9E)`, and `_info(9E)` entry points to initialize the module, as well as `dev_ops(9S)` and `cb_ops(9S)` structures.

The main interface with MAC is through a series of callbacks defined in a `mac_callbacks(9S)` structure. These callbacks control all the aspects of the device. They range from sending data, getting and setting of properties, controlling mac address filters, and also managing promiscuous mode.

The MAC framework takes care of many aspects of the device driver's management. A device that uses the MAC framework does not have to worry about creating device nodes or implementing `open(9E)` or `close(9E)` routines. In addition, all of the work to interact with `dlpi(7P)` is taken care of automatically and transparently.

**High-Level Design**

At a high-level a device driver is chiefly concerned with two different operations:

1. Sending frames
2. Receiving frames

When sending frames, the MAC framework always calls functions registered in the `mac_callbacks(9S)` structure to have the driver transmit frames on hardware. When receiving frames, the driver will generally receive an interrupt which will cause it to check for incoming data and deliver it to the MAC framework.

Configuration of a device such as whether auto-negotiation should be enabled, the speeds that the device supports, the MTU (maximum transmission unit), and the generation of pause frames are all driven by properties. The functions to get, set, and obtain information about properties are defined through callback functions specified in the `mac_callbacks(9S)` structure. The full list of properties and a description of the relevant callbacks can be found in the *PROPERTIES* section.

The MAC framework is designed to take advantage of various modern features provided by hardware, such as checksumming, segmentation offload, and hardware filtering. The MAC framework assumes none of these advanced features are present and allows device drivers to negotiate them through a capability system. Drivers can declare that they support various capabilities by implementing the optional `mc_getcapab(9E)` entry point. Each capability has its associated entry points and structures to fill out. The capabilities are detailed in the *CAPABILITIES* section.

The following sections, describe the flow of a basic device driver. For advanced device drivers, the flow is generally the same. The primary distinction is in how frames are sent and received.

### Initializing MAC Support

For a device to be used by the MAC framework, it must register with the framework and take specific actions during `_init(9E)`, `attach(9E)`, `detach(9E)`, and `_fini(9E)`.

All device drivers have to define a `dev_ops(9S)` structure which is pointed to by a `modldrv(9S)` structure and the corresponding NULL-terminated `modlinkage(9S)` structure. The `dev_ops(9S)` structure should have a `cb_ops(9S)` structure defined for it; however, it does not need to implement any of the standard `cb_ops(9S)` entry points.

Normally, in a driver's `_init(9E)` entry point, it passes its **modlinkage** structure directly to `mod_install(9F)`. To properly register with MAC, the driver must call `mac_init_ops(9F)` before it calls `mod_install(9F)`. If for some reason the `mod_install(9F)` function fails, then the driver must be removed by a call to `mac_fini_ops(9F)`.

Conversely, in the driver's `_fini(9F)` routine, it should call `mac_fini_ops(9F)` after it successfully calls `mod_remove(9F)`. For an example of how to use the `mac_init_ops(9F)` and `mac_fini_ops(9F)` functions, see the examples section in `mac_init_ops(9F)`.

### Registering with MAC

Every instance of a device should register separately with MAC. To register with MAC, a driver must allocate a `mac_register(9S)` structure, fill it in, and then call `mac_register(9F)`. The **mac\_register\_t** structure contains information about the device and all of the required function pointers that will be used as callbacks by the framework.

These steps should all be taken during a device's attach(9E) entry point. It is recommended that the driver perform this sequence of steps after the device has finished its initialization of the chipset and interrupts, though interrupts should not be enabled at that point. After it calls `mac_register(9F)` it will start receiving callbacks from the MAC framework.

To allocate the registration structure, the driver should call `mac_alloc(9F)`. Device drivers should generally always pass the symbol `MAC_VERSION` as the argument to `mac_alloc(9F)`. Upon successful completion, the driver will receive a **mac\_register\_t** structure which it should fill in. The structure and its members are documented in `mac_register(9S)`.

The `mac_callbacks(9S)` structure is not allocated as a part of the `mac_register(9S)` structure. In general, device drivers declare this statically. See the *MAC Callbacks* section for more information on how to fill it out.

Once the structure has been filled in, the driver should call `mac_register(9F)` to register itself with MAC. The handle that it uses to register with should be part of the driver's soft state. It will be used in various other support functions and callbacks.

If the call is successful, then the device driver should enable interrupts and finish any other initialization required. If the call to `mac_register(9F)` failed, then it should unwind its initialization and should return **DDI\_FAILURE** from its `attach(9E)` routine.

## MAC Callbacks

The MAC framework interacts with a device driver through a series of callbacks. These callbacks are described in their individual manual pages and the collection of callbacks is indicated in the `mac_callbacks(9S)` manual page. This section does not focus on the specific functions, but rather on interactions between them and the rest of the device driver framework.

A device driver should make no assumptions about when the various callbacks will be called and whether or not they will be called simultaneously. For example, a device driver may be asked to transmit data through a call to its `mc_tx(9F)` entry point while it is being asked to get a device property through a call to its `mc_getprop(9F)` entry point. As such, while some calls may be serialized to the device, such as setting properties, the device driver should always presume that all of its data needs to be protected with locks. While the device is holding locks, it is safe for it call the following MAC routines:

- `mac_hcksum_get(9F)`
- `mac_hcksum_set(9F)`
- `mac_lso_get(9F)`
- `mac_maxsdu_update(9F)`
- `mac_prop_info_set_default_link_flowctrl(9F)`
- `mac_prop_info_set_default_str(9F)`

- `mac_prop_info_set_default_uint8(9F)`
- `mac_prop_info_set_default_uint32(9F)`
- `mac_prop_info_set_default_uint64(9F)`
- `mac_prop_info_set_perm(9F)`
- `mac_prop_info_set_range_uint32(9F)`

Any other MAC related routines should not be called with locks held, such as `mac_link_update(9F)` or `mac_rx(9F)`. Other routines in the DDI may be called while locks are held; however, device driver writers should be careful about calling blocking routines while locks are held or in interrupt context, even when it is legal to do so.

### Receiving Data

A device driver will often receive data through the means of an interrupt or by being asked to poll for frames. When this occurs, zero or more frames, each with optional metadata, may be ready for the device driver to consume. Often each frame has a corresponding descriptor which has information about whether or not there were errors or whether or not the device successfully checksummed the packet.

During a single interrupt or poll request, a device driver should process a fixed number of frames. For each frame the device driver should:

1. Ensure that all of the DMA memory for the region is synchronized with the `ddi_dma_sync(9F)` function and the handle checked for errors if the device driver has enabled DMA error reporting as part of the Fault Management Architecture (FMA). If the driver does not rely on DMA, then it may skip this step. It is recommended that this is performed once per interrupt or poll for the entire region and not on a per-packet basis.
2. First check whether or not the frame has errors. If errors were detected, then the frame should not be sent to the operating system. It is recommended that devices keep kstats (see `kstat_create(9S)` for more information) and bump the counter whenever such an error is detected. If the device distinguishes between the types of errors, then separate kstats for each class of error are recommended. See the *STATISTICS* section for more information on the various error cases that should be considered.
3. Once the frame has been determined to be valid, the device driver should transform the frame into a `mbk(9S)`. See the section *MBLKS AND DMA* for more information on how to transform and prepare a message block.
4. If the device supports hardware checksumming (see the *CAPABILITIES* section for more information on checksumming), then the device driver should set the corresponding checksumming information with a call to `mac_hcksum_set(9F)`.

5. It should then append this new message block to the *end* of the message block chain, linking it to the **b\_next** pointer. It is vitally important that all the frames be chained in the order that they were received. If the device driver mistakenly reorders frames, then it may cause performance impacts in the TCP stack and potentially impact application correctness.

Once all the frames have been processed and assembled, the device driver should deliver them to the rest of the operating system by calling `mac_rx(9F)`. The device driver should try to give as many `mbk_t` structures to the system at once. It *should not* call `mac_rx(9F)` once for every assembled `mbk_t`.

The device driver must not hold any locks across the call to `mac_rx(9F)`. When this function is called, received data will be pushed through the networking stack and some replies may be generated and given to the driver to send out.

It is not the device driver's responsibility to determine whether or not the system can keep up with a driver's delivery rate of frames. The rest of the networking stack will handle issues related to keeping up appropriately and ensure that kernel memory is not exhausted by packets that are not being processed.

If the device driver has negotiated the `MAC_CAPAB_RINGS` capability (discussed in `mac_capab_rings(9E)`) then it should call `mac_rx_ring(9F)` and not `mac_rx(9F)`. A given interrupt may correspond to more than one ring that needs to be checked. In those cases, the driver should follow the above procedure independently for each ring. That means it will call `mac_rx_ring(9F)` once for each ring. When it is looking at the rings, the driver will need to make sure that the ring has not had interrupts disabled and is undergoing polling. This is discussed in greater detail in the `mac_capab_rings(9E)` and `mri_poll(9E)` manual pages.

Finally, the device driver should make sure that any other housekeeping activities required for the ring are taken care of such that more data can be received.

### Transmitting Data and Back Pressure

A device driver will be asked to transmit a message block chain by having its `mc_tx(9E)` entry point called. While the driver is processing the message blocks, it may run out of resources. For example, a transmit descriptor ring may become full. At that point, the device driver should return the remaining unprocessed frames. The act of returning frames indicates that the device has asserted flow control. Once this has been done, no additional calls will be made to the driver's transmit entry point and the back pressure will be propagated throughout the rest of the networking stack.

At some point in the future when resources have become available again, for example after an interrupt indicating that some portion of the transmit ring has been sent, then the device driver must notify the system that it can continue transmission. To do this, the driver should call `mac_tx_update(9F)`. After

that point, the driver will receive calls to its `mc_tx(9E)` entry point again. As mentioned in the section on callbacks, the device driver should avoid holding any particular locks across the call to `mac_tx_update(9F)`.

### **Interrupt Coalescing**

For devices operating at higher data rates, interrupt coalescing is an important part of a well functioning device and may impact the performance of the device. Not all devices support interrupt coalescing. If interrupt coalescing is supported on the device, it is recommended that device driver writers provide private properties for their device to control the interrupt coalescing rate. This will make it much easier to perform experiments and observe the impact of different interrupt rates on the rest of the system.

### **MAC Address Filter Management**

The MAC framework will attempt to use as many MAC address filters as a device has. To program a multicast address filter, the driver's `mc_multicast(9E)` entry point will be called. If the device driver runs out of filters, it should not take any special action and just return the appropriate error as documented in the corresponding manual pages for the entry points. The framework will ensure that the device is placed in promiscuous mode if it needs to.

If the hardware supports more than one unicast filters then the device driver should consider implementing the `MAC_CAPAB_RINGS` capability, which exposes a means for those multiple unicast MAC address filters to be used by the broader system. See `mac_capab_rings(9E)` for more information.

### **Receive Side Scaling**

Receive side scaling is where a hardware device supports multiple, independent queues of frames that can be received. Each of these queues is generally associated with an independent interrupt and the hardware usually performs some form of hash across the queues. Hardware which supports this, should look at implementing the `MAC_CAPAB_RINGS` capability and see `mac_capab_rings(9E)` for more information.

### **Link Updates**

It is the responsibility of the device driver to keep track of the data link's state. Many devices provide a means of receiving an interrupt when the state of the link changes. When such a change happens, the driver should update its internal data structures and then call `mac_link_update(9F)` to inform the MAC layer that this has occurred. If the device driver does not properly inform the system about link changes, then various features like link aggregations and other mechanisms that leverage the link state will not work correctly.

### **Link Speed and Auto-negotiation**

Many networking devices support more than one possible speed that they can operate at. The selection of a speed is often performed through *auto-negotiation*, though some devices allow the user to control

what speeds are advertised and used.

Logically, there are two different sets of things that the device driver needs to keep track of while it's operating:

1. The supported speeds in hardware.
2. The enabled speeds from the user.

By default, when a link first comes up, the device driver should generally configure the link to support the common set of speeds and perform auto-negotiation.

A user can control what speeds a device advertises via auto-negotiation and whether or not it performs auto-negotiation at all by using a series of properties that have **\_EN\_** in the name. These are read/write properties and there is one for each speed supported in the operating system. For a full list of them, see the *PROPERTIES* section.

In addition to these properties, there is a corresponding set of properties with **\_ADV\_** in the name. These are similar to the **\_EN\_** family of properties, but they are read-only and indicate what the device has actually negotiated. While they are generally similar to the **\_EN\_** family of properties, they may change depending on power settings. See the **Ethernet Link Properties** section in `dladm(1M)` for more information.

It's worth discussing how these different values get used throughout the different entry points. The first entry point to consider is the `mc_propinfo(9E)` entry point. For a given speed, the driver should consult whether or not the hardware supports this speed. If it does, it should fill in the default value that the hardware takes and whether or not the property is writable. The properties should also be updated to indicate whether or not it is writable. This holds for both the **\_EN\_** and **\_ADV\_** family of properties.

The next entry point is `mc_getprop(9E)`. Here, the device should first consult whether the given speed is supported. If it is not, then the driver should return `ENOTSUP`. If it does, then it should return the current value of the property.

The last property endpoint is the `mc_setprop(9E)` entry point. Here, the same logic applies. Before the driver considers whether or not the property is writable, it should first check whether or not it's a supported property. If it's not, then it should return `ENOTSUP`. Otherwise, it should proceed to check whether the property is writable, and if it is and a valid value, then it should update the property and restart the link's negotiation.

Finally, there is the `mc_getstat(9E)` entry point. Several of the statistics that are queried relate to auto-

negotiation and hardware capabilities. When a statistic relates to the hardware supporting a given speed, the **\_EN\_** properties should be ignored. The only thing that should be consulted is what the hardware itself supports. Otherwise, the statistics should look at what is currently being advertised by the device.

### Unregistering from MAC

During a driver's detach(9E) routine, it should unregister the device instance from MAC by calling `mac_unregister(9F)` on the handle that it originally called it on. If the call to `mac_unregister(9F)` failed, then the device is likely still in use and the driver should fail the call to `detach(9E)`.

### Interacting with Devices

Administrators always interact with devices through the `dladm(1M)` command line interface. The state of devices such as whether the link is considered **up** or **down**, various link properties such as the **MTU**, **auto-negotiation** state, and **flow control** state, are all exposed. It is also the preferred way that these properties are set and configured.

While device tunables may be presented in a `driver.conf(4)` file, it is recommended instead to expose such things through `dladm(1M)` private properties, whether explicitly documented or not.

## CAPABILITIES

Capabilities in the MAC Framework are optional features that a device supports which indicate various hardware features that the device supports. The two current capabilities that the system supports are related to being able to hardware perform large send offloads (LSO), often also known as TCP segmentation and the ability for hardware to calculate and verify the checksums present in IPv4, IPV6, and protocol headers such as TCP and UDP.

The MAC framework will query a device for support of a capability through the `mc_getcapab(9E)` function. Each capability has its own constant and may have corresponding data that goes along with it and a specific structure that the device is required to fill in. Note, the set of capabilities changes over time and there are also private capabilities in the system. Several of the capabilities are used in the implementation of the MAC framework. Others represent features that have not been stabilized and thus both API and binary compatibility for them is not guaranteed. It is important that the device driver handles unknown capabilities correctly. For more information, see `mc_getcapab(9E)`.

The following capabilities are stable and defined in the system:

### MAC\_CAPAB\_HCKSUM

The `MAC_CAPAB_HCKSUM` capability indicates to the system that the device driver supports some amount of checksumming. The specific data for this capability is a pointer to a `uint32_t`. To indicate no support for any kind of checksumming, the driver should either set this value to zero or simply return that it doesn't support the capability.



Note, the values that the driver declares in this capability indicate what it can do when it transmits data. If the driver can only verify checksums when receiving data, then it should not indicate that it supports this capability. The following set of flags may be combined through a bitwise inclusive OR:

#### HCKSUM\_INET\_PARTIAL

This indicates that the hardware can calculate a partial checksum for both IPv4 and IPv6; however, it requires the pseudo-header checksum be calculated for it. The pseudo-header checksum will be available for the `mbulk_t` when calling `mac_hcksum_get(9F)`. Note this does not imply that the hardware is capable of calculating the IPv4 header checksum. That should be indicated with the `HCKSUM_IPHDRCKSUM` flag.

#### HCKSUM\_INET\_FULL\_V4

This indicates that the hardware will fully calculate the L4 checksum for outgoing IPv4 packets and does not require a pseudo-header checksum. Note this does not imply that the hardware is capable of calculating the IPv4 header checksum. That should be indicated with the `HCKSUM_IPHDRCKSUM`.

#### HCKSUM\_INET\_FULL\_V6

This indicates that the hardware will fully calculate the L4 checksum for outgoing IPv6 packets and does not require a pseudo-header checksum.

#### HCKSUM\_IPHDRCKSUM

This indicates that the hardware supports calculating the checksum for the IPv4 header itself.

When in a driver's transmit function, the driver will be processing a single frame. It should call `mac_hcksum_get(9F)` to see what checksum flags are set on it. Note that the flags that are set on it are different from the ones described above and are documented in its manual page. These flags indicate how the driver is expected to program the hardware and what checksumming is required. Not all frames will require hardware checksumming or will ask the hardware to checksum it.

If a driver supports offloading the receive checksum and verification, it should check to see what the hardware indicated was verified. The driver should then call `mac_hcksum_set(9F)`. The flags used are different from the ones above and are discussed in detail in the `mac_hcksum_set(9F)` manual page. If there is no checksum information available or the driver does not support checksumming, then it should simply not call `mac_hcksum_set(9F)`.

Note that the checksum flags should be set on the first `mbulk_t` that makes up a given message. In other words, if multiple `mbulk_t` structures are linked together by the `b_cont` member to describe a single frame, then it should only be called on the first `mbulk_t` of that set. However, each distinct message

should have the checksum bits set on it, if applicable. In other words, each `mbulk_t` that is linked together by the `b_next` pointer may have checksum flags set.

It is recommended that device drivers provide a private property or `driver.conf(4)` property to control whether or not checksumming is enabled for both rx and tx; however, the default disposition is recommended to be enabled for both. This way if hardware bugs are found in the checksumming implementation, they can be disabled without requiring software updates. The transmit property should be checked when determining how to reply to `mc_getcapab(9E)` and the receive property should be checked in the context of the receive function.

### MAC\_CAPAB\_LSO

The `MAC_CAPAB_LSO` capability indicates that the driver supports various forms of large send offload (LSO). The private data is a pointer to a `mac_capab_lso_t` structure. At the moment, LSO support is limited to TCP inside of IPv4. This structure has the following members which are used to indicate various types of LSO support.

```
t_uscalar_t    lso_flags;
lso_basic_tcp_ivr4_t    lso_basic_tcp_ipv4;
```

The `lso_flags` member is used to indicate which members are valid and should be considered. Each flag represents a different form of LSO. The member should be set to the bitwise inclusive OR of the following values:

#### LSO\_TX\_BASIC\_TCP\_IPV4

This indicates hardware support for performing TCP segmentation offloading over IPv4. When this flag is set, the `lso_basic_tcp_ipv4` member must be filled in.

The `lso_basic_tcp_ipv4` member is a structure with the following members:

```
t_uscalar_t    lso_max
```

The `lso_max` member should be set to the maximum size of the TCP data payload that can be offloaded to the hardware.

Like with checksumming, it is recommended that driver writers provide a means for disabling the support of LSO even if it is enabled by default. This deals with the case where issues that pop up for LSO may be worked around without requiring additional driver work.

The following capabilities are still evolving in the operating system. They are documented such that

device driver writers may experiment with them; however, if such drivers are not present inside the core operating system repository, they may be subject to API and ABI breakage.

## MAC\_CAPAB\_RINGS

The **MAC\_CAPAB\_RINGS** capability is very important for implementing a high-performing device driver. Networking hardware structures the queues of packets to be sent and received into a ring. Each entry in this ring, has a descriptor, which describes the address and options for a packet which is going to be transmitted or received. While simple networking devices only have a single ring, many high-speed networking devices have support for many rings.

Rings are used for two important purposes. The first is receive side scaling (RSS), which is the ability to have the hardware hash the contents of a packet based on some of the protocol headers, and send it to one of several rings. These different rings may each have their own interrupt associated with them, allowing the card to receive traffic in parallel. Similar logic can be performed when sending traffic, to leverage multiple hardware resources, increasing capacity.

The second use of rings is to group them together and apply filtering rules. For example, if a packet matches a specific VLAN or MAC address, then it can be sent to a specific ring or a specific group of rings.

From the MAC framework's perspective, a driver can have one or more groups. A group consists of the following:

- One or more hardware rings.
- One or more MAC address or VLAN filters.

The details around how a device driver changes when rings are employed, the data structures that a driver must implement, and more are in a separate manual page. Please see `mac_capab_rings(9E)` for more information.

## PROPERTIES

Properties in the MAC framework represent aspects of a link. These include things like the link's current state and MTU. Many of the properties in the system are focused around auto-negotiation and controlling what link speeds are advertised. Information about properties is covered by three different device entry points. The `mc_propinfo(9E)` entry point obtains metadata about the property. The `mc_getprop(9E)` entry point obtains the property. The `mc_setprop(9E)` entry point updates the property to a new value.

Many of the properties listed below are read-only. Each property indicates whether it's read-only or it's

read/write. However, driver writers may not implement the ability to set all writable properties. Many of these depend on the card itself. In particular, all properties that relate to auto-negotiation and are read/write may not be updated if the hardware in question does not support toggling what link speeds are auto-negotiated. While copper Ethernet often does not have this restriction, it often exists with various fiber standards and phys.

The following properties are the subset of MAC framework properties that driver writers should be aware of and handle. While other properties exist in the system, driver writers should always return an error when a property not listed below is encountered. See `mc_getprop(9E)` and `mc_setprop(9E)` for more information on how to handle them.

### **MAC\_PROP\_DUPLEX**

Type: `link_duplex_t` | Permissions: **Read-Only**

The **MAC\_PROP\_DUPLEX** property is used to indicate whether or not the link is duplex. A duplex link may have traffic flowing in both directions at the same time. The `link_duplex_t` is an enumeration which may be set to any of the following values:

#### **LINK\_DUPLEX\_UNKNOWN**

The current state of the link is unknown. This may be because the link has not negotiated to a specific speed or it is down.

#### **LINK\_DUPLEX\_HALF**

The link is running at half duplex. Communication may travel in only one direction on the link at a given time.

#### **LINK\_DUPLEX\_FULL**

The link is running at full duplex. Communication may travel in both directions on the link simultaneously.

### **MAC\_PROP\_SPEED**

Type: `uint64_t` | Permissions: **Read-Only**

The **MAC\_PROP\_SPEED** property stores the current link speed in bits per second. A link that is running at 100 MBit/s would store the value 100000000ULL. A link that is running at 40 Gbit/s would store the value 40000000000ULL.

### **MAC\_PROP\_STATUS**

Type: `link_state_t` | Permissions: **Read-Only**

The **MAC\_PROP\_STATUS** property is used to indicate the current state of the link. It indicates whether the link is up or down. The **link\_state\_t** is an enumeration which may be set to any of the following values:

#### **LINK\_STATE\_UNKNOWN**

The current state of the link is unknown. This may be because the driver's **mc\_start(9E)** endpoint has not been called so it has not attempted to start the link.

#### **LINK\_STATE\_DOWN**

The link is down. This may be because of a negotiation problem, a cable problem, or some other device specific issue.

#### **LINK\_STATE\_UP**

The link is up. If auto-negotiation is in use, it should have completed. Traffic should be able to flow over the link, barring other issues.

### **MAC\_PROP\_AUTONEG**

Type: **uint8\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_AUTONEG** property indicates whether or not the device is currently configured to perform auto-negotiation. A value of **0** indicates that auto-negotiation is disabled. A **non-zero** value indicates that auto-negotiation is enabled. Devices should generally default to enabling auto-negotiation.

When getting this property, the device driver should return the current state. When setting this property, if the device supports operating in the requested mode, then the device driver should reset the link to negotiate to the new speed after updating any internal registers.

### **MAC\_PROP\_MTU**

Type: **uint32\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_MTU** property determines the maximum transmission unit (MTU). This indicates the maximum size packet that the device can transmit, ignoring its own headers. For an Ethernet device, this would exclude the size of the Ethernet header and any VLAN headers that would be placed. It is up to the driver to ensure that any MTU values that it accepts when adding in its margin and header sizes does not exceed its maximum frame size.

By default, drivers for Ethernet should initialize this value and the MTU to **1500**. When getting this property, the driver should return its current recorded MTU. When setting this property, the driver should first validate that it is within the device's valid range and then it must call

mac\_maxsdu\_update(9F). Note that the call may fail. If the call completes successfully, the driver should update the hardware with the new value of the MTU and perform any other work needed to handle it.

If the device does not support changing the MTU after the device's mc\_start(9E) entry point has been called, then driver writers should return EBUSY.

## **MAC\_PROP\_FLOWCTRL**

Type: **link\_flowctrl\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_FLOWCTRL** property manages the configuration of pause frames as part of Ethernet flow control. Note, this only describes what this device will advertise. What is actually enabled may be different and is subject to the rules of auto-negotiation. The **link\_flowctrl\_t** is an enumeration that may be set to one of the following values:

### **LINK\_FLOWCTRL\_NONE**

Flow control is disabled. No pause frames should be generated or honored.

### **LINK\_FLOWCTRL\_RX**

The device can receive pause frames; however, it should not generate them.

### **LINK\_FLOWCTRL\_TX**

The device can generate pause frames; however, it does not support receiving them.

### **LINK\_FLOWCTRL\_BI**

The device supports both sending and receiving pause frames.

When getting this property, the device driver should return the way that it has configured the device, not what the device has actually negotiated. When setting the property, it should update the hardware and allow the link to potentially perform auto-negotiation again.

The remaining properties are all about various auto-negotiation link speeds. They fall into two different buckets: properties with **\_ADV\_** in the name and properties with **\_EN\_** in the name. For any given supported speed, there is one of each. The **\_EN\_** set of properties are read/write properties that control what should be advertised by the device. When these are retrieved, they should return the current value of the property. When they are set, they should change how the hardware advertises the specific speed and trigger any kind of link reset and auto-negotiation, if enabled, to occur.

The **\_ADV\_** set of properties are read-only properties. They are meant to reflect what has actually been negotiated. These may be different from the **\_EN\_** family of properties, especially when different power

management settings are at play.

See the *Link Speed and Auto-negotiation* section for more information.

The properties are ordered in increasing link speed:

#### **MAC\_PROP\_ADV\_10HDX\_CAP**

Type: **uint8\_t** | Permissions: **Read-Only**

The **MAC\_PROP\_ADV\_10HDX\_CAP** property describes whether or not 10 Mbit/s half-duplex support is advertised.

#### **MAC\_PROP\_EN\_10HDX\_CAP**

Type: **uint8\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_EN\_10HDX\_CAP** property describes whether or not 10 Mbit/s half-duplex support is enabled.

#### **MAC\_PROP\_ADV\_10FDX\_CAP**

Type: **uint8\_t** | Permissions: **Read-Only**

The **MAC\_PROP\_ADV\_10FDX\_CAP** property describes whether or not 10 Mbit/s full-duplex support is advertised.

#### **MAC\_PROP\_EN\_10FDX\_CAP**

Type: **uint8\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_EN\_10FDX\_CAP** property describes whether or not 10 Mbit/s full-duplex support is enabled.

#### **MAC\_PROP\_ADV\_100HDX\_CAP**

Type: **uint8\_t** | Permissions: **Read-Only**

The **MAC\_PROP\_ADV\_100HDX\_CAP** property describes whether or not 100 Mbit/s half-duplex support is advertised.

#### **MAC\_PROP\_EN\_100HDX\_CAP**

Type: **uint8\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_EN\_100HDX\_CAP** property describes whether or not 100 Mbit/s half-

duplex support is enabled.

#### **MAC\_PROP\_ADV\_100FDX\_CAP**

Type: **uint8\_t** | Permissions: **Read-Only**

The **MAC\_PROP\_ADV\_100FDX\_CAP** property describes whether or not 100 Mbit/s full-duplex support is advertised.

#### **MAC\_PROP\_EN\_100FDX\_CAP**

Type: **uint8\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_EN\_100FDX\_CAP** property describes whether or not 100 Mbit/s full-duplex support is enabled.

#### **MAC\_PROP\_ADV\_100T4\_CAP**

Type: **uint8\_t** | Permissions: **Read-Only**

The **MAC\_PROP\_ADV\_100T4\_CAP** property describes whether or not 100 Mbit/s Ethernet using the 100BASE-T4 standard is advertised.

#### **MAC\_PROP\_EN\_100T4\_CAP**

Type: **uint8\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_ADV\_100T4\_CAP** property describes whether or not 100 Mbit/s Ethernet using the 100BASE-T4 standard is enabled.

#### **MAC\_PROP\_ADV\_1000HDX\_CAP**

Type: **uint8\_t** | Permissions: **Read-Only**

The **MAC\_PROP\_ADV\_1000HDX\_CAP** property describes whether or not 1 Gbit/s half-duplex support is advertised.

#### **MAC\_PROP\_EN\_1000HDX\_CAP**

Type: **uint8\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_EN\_1000HDX\_CAP** property describes whether or not 1 Gbit/s half-duplex support is enabled.

#### **MAC\_PROP\_ADV\_1000FDX\_CAP**

Type: **uint8\_t** | Permissions: **Read-Only**



The **MAC\_PROP\_ADV\_1000FDX\_CAP** property describes whether or not 1 Gbit/s full-duplex support is advertised.

#### **MAC\_PROP\_EN\_1000FDX\_CAP**

Type: **uint8\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_EN\_1000FDX\_CAP** property describes whether or not 1 Gbit/s full-duplex support is enabled.

#### **MAC\_PROP\_ADV\_2500FDX\_CAP**

Type: **uint8\_t** | Permissions: **Read-Only**

The **MAC\_PROP\_ADV\_2500FDX\_CAP** property describes whether or not 2.5 Gbit/s full-duplex support is advertised.

#### **MAC\_PROP\_EN\_2500FDX\_CAP**

Type: **uint8\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_EN\_2500FDX\_CAP** property describes whether or not 2.5 Gbit/s full-duplex support is enabled.

#### **MAC\_PROP\_ADV\_5000FDX\_CAP**

Type: **uint8\_t** | Permissions: **Read-Only**

The **MAC\_PROP\_ADV\_5000FDX\_CAP** property describes whether or not 5.0 Gbit/s full-duplex support is advertised.

#### **MAC\_PROP\_EN\_5000FDX\_CAP**

Type: **uint8\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_EN\_5000FDX\_CAP** property describes whether or not 5.0 Gbit/s full-duplex support is enabled.

#### **MAC\_PROP\_ADV\_10GFDX\_CAP**

Type: **uint8\_t** | Permissions: **Read-Only**

The **MAC\_PROP\_ADV\_10GFDX\_CAP** property describes whether or not 10 Gbit/s full-duplex support is advertised.

#### **MAC\_PROP\_EN\_10GFDX\_CAP**

Type: **uint8\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_EN\_10GFDX\_CAP** property describes whether or not 10 Gbit/s full-duplex support is enabled.

#### **MAC\_PROP\_ADV\_40GFDX\_CAP**

Type: **uint8\_t** | Permissions: **Read-Only**

The **MAC\_PROP\_ADV\_40GFDX\_CAP** property describes whether or not 40 Gbit/s full-duplex support is advertised.

#### **MAC\_PROP\_EN\_40GFDX\_CAP**

Type: **uint8\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_EN\_40GFDX\_CAP** property describes whether or not 40 Gbit/s full-duplex support is enabled.

#### **MAC\_PROP\_ADV\_100GFDX\_CAP**

Type: **uint8\_t** | Permissions: **Read-Only**

The **MAC\_PROP\_ADV\_100GFDX\_CAP** property describes whether or not 100 Gbit/s full-duplex support is advertised.

#### **MAC\_PROP\_EN\_100GFDX\_CAP**

Type: **uint8\_t** | Permissions: **Read/Write**

The **MAC\_PROP\_EN\_100GFDX\_CAP** property describes whether or not 100 Gbit/s full-duplex support is enabled.

### **Private Properties**

In addition to the defined properties above, drivers are allowed to define private properties. These private properties are device-specific properties. All private properties share the same constant, **MAC\_PROP\_PRIVATE**. Properties are distinguished by a name, which is a character string. The list of such private properties is defined when registering with mac in the **m\_priv\_props** member of the `mac_register(9S)` structure.

The driver may define whatever semantics it wants for these private properties. They will not be listed when running `dladm(1M)`, unless explicitly requested by name. All such properties should start with a leading underscore character and then consist of alphanumeric ASCII characters and additional underscores or hyphens.

Properties of type **MAC\_PROP\_PRIVATE** may show up in all three property related entry points: `mc_propinfo(9E)`, `mc_getprop(9E)`, and `mc_setprop(9E)`. Device drivers should tell the different properties apart by using the `strcmp(9F)` function to compare it to the set of properties that it knows about. When encountering properties that it doesn't know, it should treat them like all other unknown properties.

## STATISTICS

The MAC framework defines a couple different sets of statistics which are based on various standards for devices to implement. Statistics are retrieved through the `mc_getstat(9E)` entry point. There are both statistics that are required for all devices and then there is a separate set of Ethernet specific statistics. Not all devices will support every statistic. In many cases, several device registers will need to be combined to create the proper stat.

In general, if the device is not keeping track of these statistics, then it is recommended that the driver store these values as a **uint64\_t** to ensure that overflow does not occur.

If a device does not support a specific statistic, then it is fine to return that it is not supported. The same should be used for unrecognized statistics. See `mc_getstat(9E)` for more information on the proper way to handle these.

### General Device Statistics

The following statistics are based on MIB-II statistics from both RFC 1213 and RFC 1573.

#### **MAC\_STAT\_IFSPEED**

The device's current speed in bits per second.

#### **MAC\_STAT\_MULTIRCV**

The total number of received multicast packets.

#### **MAC\_STAT\_BRDCSTRCV**

The total number of received broadcast packets.

#### **MAC\_STAT\_MULTIXMT**

The total number of transmitted multicast packets.

#### **MAC\_STAT\_BRDCSTXMT**

The total number of received broadcast packets.

#### **MAC\_STAT\_NORCVBUF**

The total number of packets discarded by the hardware due to a lack of receive buffers.

**MAC\_STAT\_IERRORS**

The total number of errors detected on input.

**MAC\_STAT\_UNKNOWNNS**

The total number of received packets that were discarded because they were of an unknown protocol.

**MAC\_STAT\_NOXMTBUF**

The total number of outgoing packets dropped due to a lack of transmit buffers.

**MAC\_STAT\_OERRORS**

The total number of outgoing packets that resulted in errors.

**MAC\_STAT\_COLLISIONS**

Total number of collisions encountered by the transmitter.

**MAC\_STAT\_RBYTES**

The total number of **bytes** received by the device, regardless of packet type.

**MAC\_STAT\_IPACKETS**

The total number of **packets** received by the device, regardless of packet type.

**MAC\_STAT\_OBYTES**

The total number of **bytes** transmitted by the device, regardless of packet type.

**MAC\_STAT\_OPACKETS**

The total number of **packets** sent by the device, regardless of packet type.

**MAC\_STAT\_UNDERFLOWS**

The total number of packets that were smaller than the minimum sized packet for the device and were therefore dropped.

**MAC\_STAT\_OVERFLOWS**

The total number of packets that were larger than the maximum sized packet for the device and were therefore dropped.

**Ethernet Specific Statistics**

The following statistics are specific to Ethernet devices. They refer to values from RFC 1643 and include various MII/GMII specific stats. Many of these are also defined in IEEE 802.3.

**ETHER\_STAT\_ADV\_CAP\_1000FDX**

Indicates that the device is advertising support for 1 Gbit/s full-duplex operation.

**ETHER\_STAT\_ADV\_CAP\_1000HDX**

Indicates that the device is advertising support for 1 Gbit/s half-duplex operation.

**ETHER\_STAT\_ADV\_CAP\_100FDX**

Indicates that the device is advertising support for 100 Mbit/s full-duplex operation.

**ETHER\_STAT\_ADV\_CAP\_100GFDX**

Indicates that the device is advertising support for 100 Gbit/s full-duplex operation.

**ETHER\_STAT\_ADV\_CAP\_100HDX**

Indicates that the device is advertising support for 100 Mbit/s half-duplex operation.

**ETHER\_STAT\_ADV\_CAP\_100T4**

Indicates that the device is advertising support for 100 Mbit/s 100BASE-T4 operation.

**ETHER\_STAT\_ADV\_CAP\_10FDX**

Indicates that the device is advertising support for 10 Mbit/s full-duplex operation.

**ETHER\_STAT\_ADV\_CAP\_10GFDX**

Indicates that the device is advertising support for 10 Gbit/s full-duplex operation.

**ETHER\_STAT\_ADV\_CAP\_10HDX**

Indicates that the device is advertising support for 10 Mbit/s half-duplex operation.

**ETHER\_STAT\_ADV\_CAP\_2500FDX**

Indicates that the device is advertising support for 2.5 Gbit/s full-duplex operation.

**ETHER\_STAT\_ADV\_CAP\_40GFDX**

Indicates that the device is advertising support for 40 Gbit/s full-duplex operation.

**ETHER\_STAT\_ADV\_CAP\_5000FDX**

Indicates that the device is advertising support for 5.0 Gbit/s full-duplex operation.

**ETHER\_STAT\_ADV\_CAP\_ASMPAUSE**

Indicates that the device is advertising support for receiving pause frames.

**ETHER\_STAT\_ADV\_CAP\_AUTONEG**

Indicates that the device is advertising support for auto-negotiation.

**ETHER\_STAT\_ADV\_CAP\_PAUSE**

Indicates that the device is advertising support for generating pause frames.

**ETHER\_STAT\_ADV\_REMFAULT**

Indicates that the device is advertising support for detecting faults in the remote link peer.

**ETHER\_STAT\_ALIGN\_ERRORS**

Indicates the number of times an alignment error was generated by the Ethernet device. This is a count of packets that were not an integral number of octets and failed the FCS check.

**ETHER\_STAT\_CAP\_1000FDX**

Indicates the device supports 1 Gbit/s full-duplex operation.

**ETHER\_STAT\_CAP\_1000HDX**

Indicates the device supports 1 Gbit/s half-duplex operation.

**ETHER\_STAT\_CAP\_100FDX**

Indicates the device supports 100 Mbit/s full-duplex operation.

**ETHER\_STAT\_CAP\_100GFDX**

Indicates the device supports 100 Gbit/s full-duplex operation.

**ETHER\_STAT\_CAP\_100HDX**

Indicates the device supports 100 Mbit/s half-duplex operation.

**ETHER\_STAT\_CAP\_100T4**

Indicates the device supports 100 Mbit/s 100BASE-T4 operation.

**ETHER\_STAT\_CAP\_10FDX**

Indicates the device supports 10 Mbit/s full-duplex operation.

**ETHER\_STAT\_CAP\_10GFDX**

Indicates the device supports 10 Gbit/s full-duplex operation.

**ETHER\_STAT\_CAP\_10HDX**

Indicates the device supports 10 Mbit/s half-duplex operation.

**ETHER\_STAT\_CAP\_2500FDX**

Indicates the device supports 2.5 Gbit/s full-duplex operation.

**ETHER\_STAT\_CAP\_40GFDX**

Indicates the device supports 40 Gbit/s full-duplex operation.

**ETHER\_STAT\_CAP\_5000FDX**

Indicates the device supports 5.0 Gbit/s full-duplex operation.

**ETHER\_STAT\_CAP\_ASMPAUSE**

Indicates that the device supports the ability to receive pause frames.

**ETHER\_STAT\_CAP\_AUTONEG**

Indicates that the device supports the ability to perform link auto-negotiation.

**ETHER\_STAT\_CAP\_PAUSE**

Indicates that the device supports the ability to transmit pause frames.

**ETHER\_STAT\_CAP\_REMFAULT**

Indicates that the device supports the ability of detecting a remote fault in a link peer.

**ETHER\_STAT\_CARRIER\_ERRORS**

Indicates the number of times that the Ethernet carrier sense condition was lost or not asserted.

**ETHER\_STAT\_DEFER\_XMTS**

Indicates the number of frames for which the device was unable to transmit the frame due to being busy and had to try again.

**ETHER\_STAT\_EX\_COLLISIONS**

Indicates the number of frames that failed to send due to an excessive number of collisions.

**ETHER\_STAT\_FCS\_ERRORS**

Indicates the number of times that a frame check sequence failed.

**ETHER\_STAT\_FIRST\_COLLISIONS**

Indicates the number of times that a frame was eventually transmitted successfully, but only after a single collision.

**ETHER\_STAT\_JABBER\_ERRORS**

Indicates the number of frames that were received that were both larger than the maximum packet size and failed the frame check sequence.

**ETHER\_STAT\_LINK\_ASMPAUSE**

Indicates whether the link is currently configured to accept pause frames.

**ETHER\_STAT\_LINK\_AUTONEG**

Indicates whether the current link state is a result of auto-negotiation.

**ETHER\_STAT\_LINK\_DUPLEX**

Indicates the current duplex state of the link. The values used here should be the same as documented for **MAC\_PROP\_DUPLEX**.

**ETHER\_STAT\_LINK\_PAUSE**

Indicates whether the link is currently configured to generate pause frames.

**ETHER\_STAT\_LP\_CAP\_1000FDX**

Indicates the remote device supports 1 Gbit/s full-duplex operation.

**ETHER\_STAT\_LP\_CAP\_1000HDX**

Indicates the remote device supports 1 Gbit/s half-duplex operation.

**ETHER\_STAT\_LP\_CAP\_100FDX**

Indicates the remote device supports 100 Mbit/s full-duplex operation.

**ETHER\_STAT\_LP\_CAP\_100GFDX**

Indicates the remote device supports 100 Gbit/s full-duplex operation.

**ETHER\_STAT\_LP\_CAP\_100HDX**

Indicates the remote device supports 100 Mbit/s half-duplex operation.

**ETHER\_STAT\_LP\_CAP\_100T4**

Indicates the remote device supports 100 Mbit/s 100BASE-T4 operation.

**ETHER\_STAT\_LP\_CAP\_10FDX**

Indicates the remote device supports 10 Mbit/s full-duplex operation.

**ETHER\_STAT\_LP\_CAP\_10GFDX**

Indicates the remote device supports 10 Gbit/s full-duplex operation.

**ETHER\_STAT\_LP\_CAP\_10HDX**

Indicates the remote device supports 10 Mbit/s half-duplex operation.



**ETHER\_STAT\_LP\_CAP\_2500FDX**

Indicates the remote device supports 2.5 Gbit/s full-duplex operation.

**ETHER\_STAT\_LP\_CAP\_40GFDX**

Indicates the remote device supports 40 Gbit/s full-duplex operation.

**ETHER\_STAT\_LP\_CAP\_5000FDX**

Indicates the remote device supports 5.0 Gbit/s full-duplex operation.

**ETHER\_STAT\_LP\_CAP\_ASMPAUSE**

Indicates that the remote device supports the ability to receive pause frames.

**ETHER\_STAT\_LP\_CAP\_AUTONEG**

Indicates that the remote device supports the ability to perform link auto-negotiation.

**ETHER\_STAT\_LP\_CAP\_PAUSE**

Indicates that the remote device supports the ability to transmit pause frames.

**ETHER\_STAT\_LP\_CAP\_REMFAULT**

Indicates that the remote device supports the ability of detecting a remote fault in a link peer.

**ETHER\_STAT\_MACRCV\_ERRORS**

Indicates the number of times that the internal MAC layer encountered an error when attempting to receive and process a frame.

**ETHER\_STAT\_MACXMT\_ERRORS**

Indicates the number of times that the internal MAC layer encountered an error when attempting to process and transmit a frame.

**ETHER\_STAT\_MULTI\_COLLISIONS**

Indicates the number of times that a frame was eventually transmitted successfully, but only after more than one collision.

**ETHER\_STAT\_SQE\_ERRORS**

Indicates the number of times that an SQE error occurred. The specific conditions for this error are documented in IEEE 802.3.

**ETHER\_STAT\_TOOLONG\_ERRORS**

Indicates the number of frames that were received that were longer than the maximum frame size supported by the device.

**ETHER\_STAT\_TOOSHORT\_ERRORS**

Indicates the number of frames that were received that were shorter than the minimum frame size supported by the device.

**ETHER\_STAT\_TX\_LATE\_COLLISIONS**

Indicates the number of times a collision was detected late on the device.

**ETHER\_STAT\_XCVR\_ADDR**

Indicates the address of the MII/GMII receiver address.

**ETHER\_STAT\_XCVR\_ID**

Indicates the id of the MII/GMII receiver address.

**ETHER\_STAT\_XCVR\_INUSE**

Indicates what kind of receiver is in use. The following values may be used:

**XCVR\_UNDEFINED**

The receiver type is undefined by the hardware.

**XCVR\_NONE**

There is no receiver in use by the hardware.

**XCVR\_10**

The receiver supports 10BASE-T operation.

**XCVR\_100T4**

The receiver supports 100BASE-T4 operation.

**XCVR\_100X**

The receiver supports 100BASE-TX operation.

**XCVR\_100T2**

The receiver supports 100BASE-T2 operation.

**XCVR\_1000X**

The receiver supports 1000BASE-X operation. This is used for all fiber receivers.

**XCVR\_1000T**

The receiver supports 1000BASE-T operation. This is used for all copper receivers.

**Device Specific kstats**

In addition to the defined statistics above, if the device driver maintains additional statistics or the device provides additional statistics, it should create its own kstats through the `kstat_create(9F)` function to allow operators to observe them.

**TX STALL DETECTION, DEVICE RESETS, AND FAULT MANAGEMENT**

Device drivers are the first line of defense for dealing with broken devices and bugs in their firmware. While most devices will rarely fail, it is important that when designing and implementing the device driver that particular attention is paid in the design with respect to RAS (Reliability, Availability, and Serviceability). While everything described in this section is optional, it is highly recommended that all new device drivers follow these guidelines.

The Fault Management Architecture (FMA) provides facilities for detecting and reporting various classes of defects and faults. Specifically for networking device drivers, issues that should be detected and reported include:

- ⊕ Device internal uncorrectable errors
- ⊕ Device internal correctable errors
- ⊕ PCI and PCI Express transport errors
- ⊕ Device temperature alarms
- ⊕ Device transmission stalls
- ⊕ Device communication timeouts
- ⊕ High invalid interrupts

All such errors fall into three primary categories:

1. Errors detected by the Fault Management Architecture
2. Errors detected by the device and indicated to the device driver
3. Errors detected by the device driver

**Fault Management Setup and Teardown**

Drivers should initialize support for the fault management framework by calling `ddi_fm_init(9F)` from

their attach(9E) routine. By registering with the fault management framework, a device driver is given the chance to detect and notice transport errors as well as report other errors that exist. While a device driver does not need to indicate that it is capable of all such capabilities described in `ddi_fm_init(9F)`, we suggest that device drivers at least register the **DDI\_FM\_EREPORT\_CAPABLE** so as to allow the driver to report issues that it detects.

If the driver registers with the fault management framework during its attach(9E) entry point, it must call `ddi_fm_fini(9E)` during its detach(9E) entry point.

### Transport Errors

Many modern networking devices leverage PCI or PCI Express. As such, there are two primary ways that device drivers access data: they either memory map device registers and use routines like `ddi_get8(9F)` and `ddi_put8(9F)` or they use direct memory access (DMA). New device drivers should always enable checking of the transport layer by marking their support in the `ddi_device_acc_attr_t(9S)` structure and using routines like `ddi_fm_acc_err_get(9F)` and `ddi_fm_dma_err_get(9F)` to detect if errors have occurred.

### Device Indicated Errors

Many devices have capabilities to announce to a device driver that a fatal correctable error or uncorrectable error has occurred. Other devices have the ability to indicate that various physical issues have occurred such as a fan failing or a temperature sensor having fired.

Drivers should wire themselves to receive notifications when these events occur. The means and capabilities will vary from device to device. For example, some devices will generate information about these notifications through special interrupts. Other devices may have a register that software can poll. In the cases where polling is required, driver writers should try not to poll too frequently and should generally only poll when the device is actively being used, e.g. between calls to the `mc_start(9E)` and `mc_stop(9E)` entry points.

### Driver Transmit Stall Detection

One of the primary responsibilities of a hardened device driver is to perform transmit stall detection. The core idea behind tx stall detection is that the driver should record when it's getting activity related to when data has been successfully transmitted. Most devices should be transmitting data on a regular basis as long as the link is up. If it is not, then this may indicate that the device is stuck and needs to be reset. At this time, the MAC framework does not provide any resources for performing these checks; however, polling on each individual transmit ring for the last completion time while something is actively being transmitted through the use of routines such as `timeout(9F)` may be a reasonable starting point.

### Driver Command Timeout Detection

Each device is programmed in different ways. Some devices are programmed through asynchronous commands while others are programmed by writing directly to memory mapped registers. If a device receives asynchronous replies to commands, then the device driver should set reasonable timeouts for all such commands and plan on detecting them. If a timeout occurs, the driver should presume that there is an issue with the hardware and proceed to abort the command or reset the device.

Many devices do not have such a communication mechanism. However, whenever there is some activity where the device driver must wait, then it should be prepared for the fact that the device may never get back to it and react appropriately by performing some kind of device reset.

### Reacting to Errors

When any of the above categories of errors has been triggered, the behavior that the device driver should take depends on the kind of error. If a fatal error, for example, a transport error, a transmit stall was detected, or the device indicated an uncorrectable error was detected, then it is important that the driver take the following steps:

1. Set a flag in the device driver's state that indicates that it has hit an error condition. When this error condition flag is asserted, transmitted packets should be accepted and dropped and actions that would require writing to the device state should fail with an error. This flag should remain until the device has been successfully restarted.
2. If the error was not a transport error that was indicated by the fault management architecture, e.g. a transport error that was detected, then the device driver should post an **ereport** indicating what has occurred with the `ddi_fm_ereport_post(9F)` function.
3. The device driver should indicate that the device's service was lost with a call to `ddi_fm_service_impact(9F)` using the symbol **DDI\_SERVICE\_LOST**.
4. At this point the device driver should issue a device reset through some device-specific means.
5. When the device reset has been completed, then the device driver should restore all of the programmed state to the device. This includes things like the current MTU, advertised auto-negotiation speeds, MAC address filters, and more.
6. Finally, when service has been restored, the device driver should call `ddi_fm_service_impact(9F)` using the symbol **DDI\_SERVICE\_RESTORED**.

When a non-fatal error occurs, then the device driver should submit an ereport and should optionally mark the device degraded using `ddi_fm_service_impact(9F)` with the **DDI\_SERVICE\_DEGRADED**

value depending on the nature of the problem that has occurred.

Device drivers should never make the decision to remove a device from service based on errors that have occurred nor should they panic the system. Rather, the device driver should always try to notify the operating system with various ereports and allow its policy decisions to occur. The decision to retire a device lies in the hands of the fault management architecture. It knows more about the operator's intent and the surrounding system's state than the device driver itself does and it will make the call to offline and retire the device if it is required.

### Device Resets

When resetting a device, a device driver must exercise caution. If a device driver has not been written to plan for a device reset, then it may not correctly restore the device's state after such a reset. Such state should be stored in the instance's private state data as the MAC framework does not know about device resets and will not inform the device again about the expected, programmed state.

One wrinkle with device resets is that many networking cards show up as multiple PCI functions on a single device, for example, each port may show up as a separate function and thus have a separate instance of the device driver attached. When resetting a function, device driver writers should carefully read the device programming manuals and verify whether or not a reset impacts only the stalled function or if it impacts all function across the device.

If the only way to reset a given function is through the device, then this may require more coordination and work on the part of the device driver to ensure that all the other instances are correctly restored. In cases where this occurs, some devices offer ways of injecting interrupts onto those other functions to notify them that this is occurring.

### MBLKS AND DMA

The networking stack manages framed data through the use of the mblk(9S) structure. The mblk allows for a single message to be made up of individual blocks. Each part is linked together through its **b\_cont** member. However, it also allows for multiple messages to be chained together through the use of the **b\_next** member. While the networking stack works with these structures, device drivers generally work with DMA regions. There are two different strategies that device drivers use for handling these two different cases: copying and binding.

### Copying Data

The first way that device drivers handle interfacing between the two is by having two separate regions of memory. One part is memory which has been allocated for DMA through a call to `ddi_dma_alloc(9F)` and the other is memory associated with the memory block.

In this case, a driver will use `bcopy(9F)` to copy memory between the two distinct regions. When

transmitting a packet, it will copy the memory from the `mblk_t` to the DMA region. When receiving memory, it will allocate a `mblk_t` through the `allocb(9F)` routine, copy the memory across with `bcopy(9F)`, and then increment the `mblk_t`'s `w_ptr` structure.

If, when receiving, memory is not available for a new message block, then the frame should be skipped and effectively dropped. A `kstat` should be bumped when such an occasion occurs.

### **Binding Data**

An alternative approach to copying data is to use DMA binding. When using DMA binding, the OS takes care of mapping between DMA memory and normal device memory. The exact process is a bit different between transmit and receive.

When transmitting a device driver has an `mblk_t` and needs to call the `ddi_dma_addr_bind_handle(9F)` function to bind it to an already existing DMA handle. At that point, it will receive various DMA cookies that it can use to obtain the addresses to program the device with for transmitting data. Once the transmit is done, the driver must then make sure to call `freemsg(9F)` to release the data. It must not call `freemsg(9F)` before it receives an interrupt from the device indicating that the data has been transmitted, otherwise it risks sending arbitrary kernel memory.

When receiving data, the device can perform a similar operation. First, it must bind the DMA memory into the kernel's virtual memory address space through a call to the `ddi_dma_addr_bind_handle(9F)` function if it has not already. Once it has, it must then call `desballoc(9F)` to try and create a new `mblk_t` which leverages the associated memory. It can then pass that `mblk_t` up to the stack.

### **Considerations**

When deciding which of these options to use, there are many different considerations that must be made. The answer as to whether to bind memory or to copy data is not always simpler.

The first thing to remember is that DMA resources may be finite on a given platform. Consider the case of receiving data. A device driver that binds one of its receive descriptors may not get it back for quite some time as it may be used by the kernel until an application actually consumes it. Device drivers that try to bind memory for receive, often work with the constraint that they must be able to replace that DMA memory with another DMA descriptor. If they were not replaced, then eventually the device would not be able to receive additional data into the ring.

On the other hand, particularly for larger frames, copying every packet from one buffer to another can be a source of additional latency and memory waste in the system. For larger copies, the cost of copying may dwarf any potential cost of performing DMA binding.

For device driver authors that are unsure of what to do, they should first employ the copying method to

simplify the act of writing the device driver. The copying method is simpler and also allows the device driver author not to worry about allocated DMA memory that is still outstanding when it is asked to unload.

If device driver writers are worried about the cost, it is recommended to make the decision as to whether or not to copy or bind DMA data a separate private property for both transmitting and receiving. That private property should indicate the size of the received frame at which to switch from one format to the other. This way, data can be gathered to determine what the impact of each method is on a given platform.

## SEE ALSO

dladm(1M), driver.conf(4), ieee802.3(5), dlpi(7P), \_fini(9E), \_info(9E), \_init(9E), attach(9E), close(9E), detach(9E), mac\_capab\_rings(9E), mc\_close(9E), mc\_getcapab(9E), mc\_getprop(9E), mc\_getstat(9E), mc\_multicast(9E), mc\_open(9E), mc\_propinfo(9E), mc\_setpromisc(9E), mc\_setprop(9E), mc\_start(9E), mc\_stop(9E), mc\_tx(9E), mc\_unicst(9E), open(9E), allocb(9F), bcopy(9F), ddi\_dma\_addr\_bind\_handle(9F), ddi\_dma\_alloc(9F), ddi\_fm\_acc\_err\_get(9F), ddi\_fm\_dma\_err\_get(9F), ddi\_fm\_ereport\_post(9F), ddi\_fm\_fini(9F), ddi\_fm\_init(9F), ddi\_fm\_service\_impact(9F), ddi\_get8(9F), ddi\_put8(9F), desballoc(9F), freemsg(9F), kstat\_create(9F), mac\_alloc(9F), mac\_fini\_ops(9F), mac\_hcksum\_get(9F), mac\_hcksum\_set(9F), mac\_init\_ops(9F), mac\_link\_update(9F), mac\_lso\_get(9F), mac\_maxsdu\_update(9F), mac\_prop\_info\_set\_default\_link\_flowctrl(9F), mac\_prop\_info\_set\_default\_str(9F), mac\_prop\_info\_set\_default\_uint32(9F), mac\_prop\_info\_set\_default\_uint64(9F), mac\_prop\_info\_set\_default\_uint8(9F), mac\_prop\_info\_set\_perm(9F), mac\_prop\_info\_set\_range\_uint32(9F), mac\_register(9F), mac\_rx(9F), mac\_unregister(9F), mc\_getprop(9F), mc\_tx(9F), mod\_install(9F), mod\_remove(9F), strcmp(9F), timeout(9F), cb\_ops(9S), ddi\_device\_acc\_attr\_t(9S), dev\_ops(9S), kstat\_create(9S), mac\_callbacks(9S), mac\_register(9S), mblk(9S), modldrv(9S), modlinkage(9S)

McCloghrie, K. and Rose, M., *RFC 1213 Management Information Base for Network Management of TCP/IP-based internets: MIB-II*, March 1991.

McCloghrie, K. and Kastenholtz, F., *RFC 1573 Evolution of the Interfaces Group of MIB-II*, January 1994.

Kastenholtz, F., *RFC 1643 Definitions of Managed Objects for the Ethernet-like, Interface Types*.